

# If It's Not Deterministic, It's Crap: Deterministic Machine Learning and Molecular Dynamics

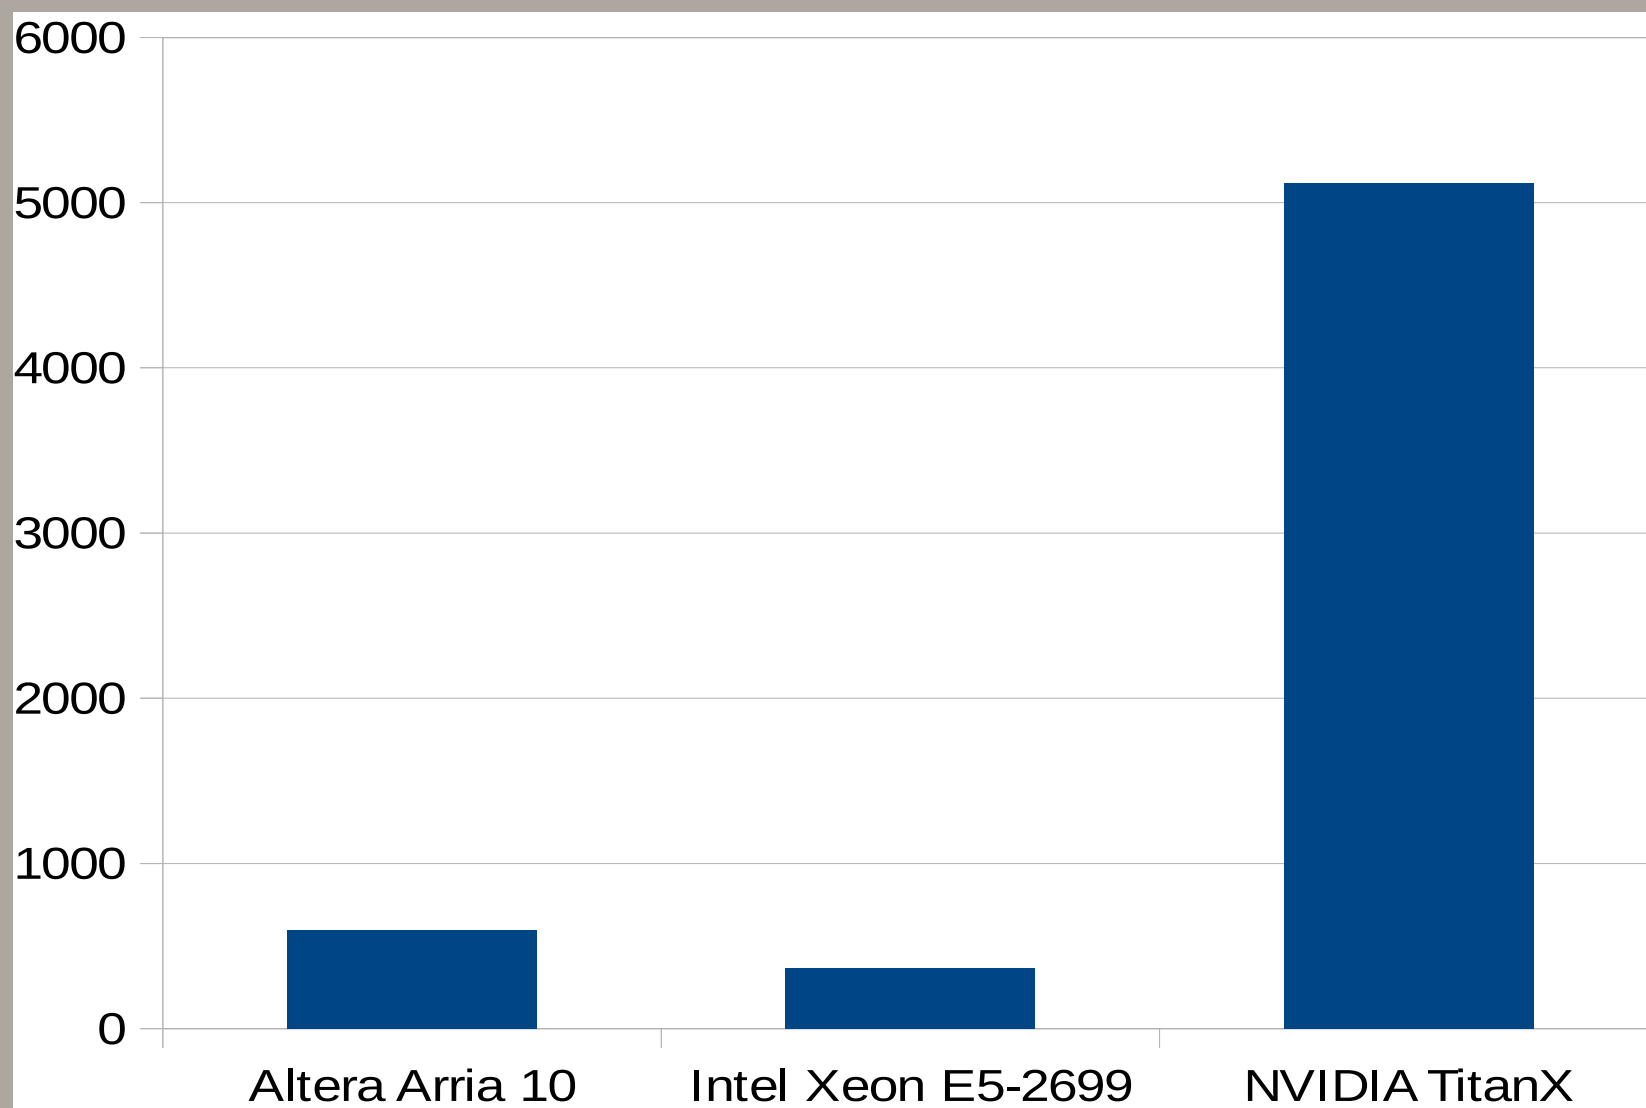
# Spoilers

- GPUs/FPGAs/CPUs/ASICs ad nauseum
- AMBER Molecular Dynamics
- Determinism Matters
- Multi-GPU Servers
- Neural Networks
- Deterministic Model Parallelism
- DGX1: \$129K of Aspirational Computing for the 1%

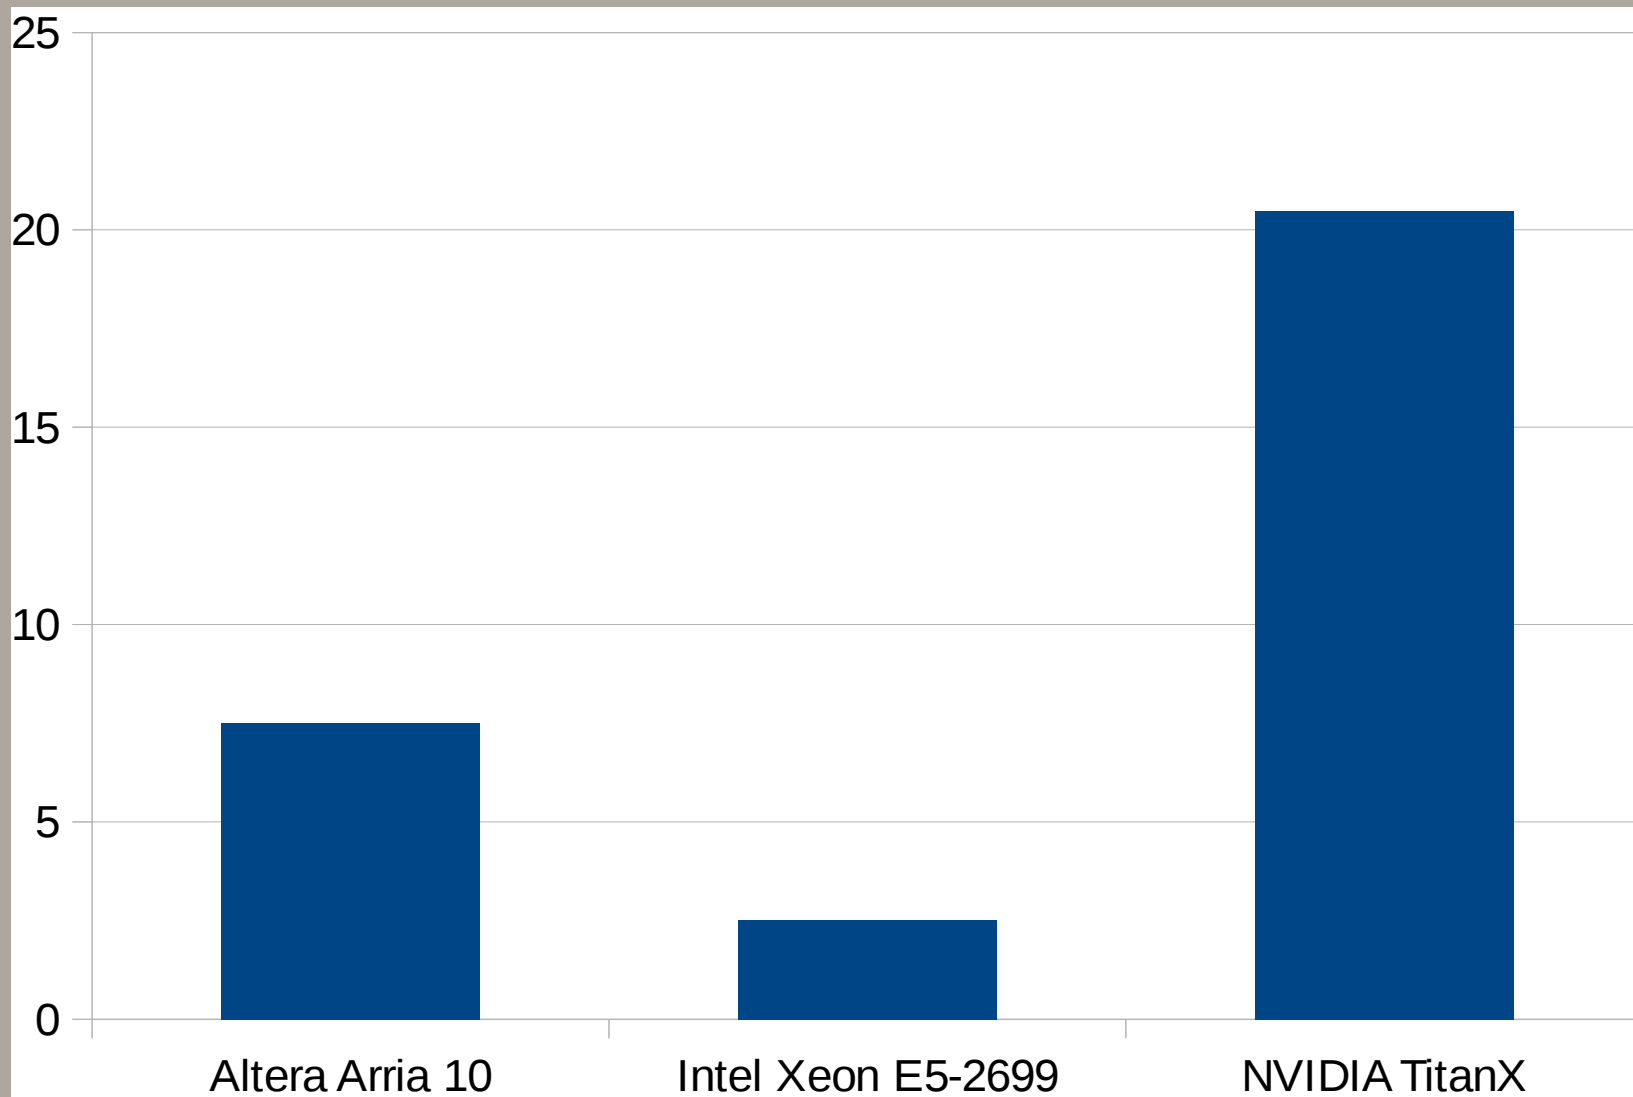
# 2016 TLDR: It's (still) the GPUs, Stupid

- Despite new hardware from Altera, IBM and Intel, not much has changed
- Intel/Altera training performance sucks
- Intel/Altera prediction performance also sucks (just not quite as much)

# AlexNet Images/s

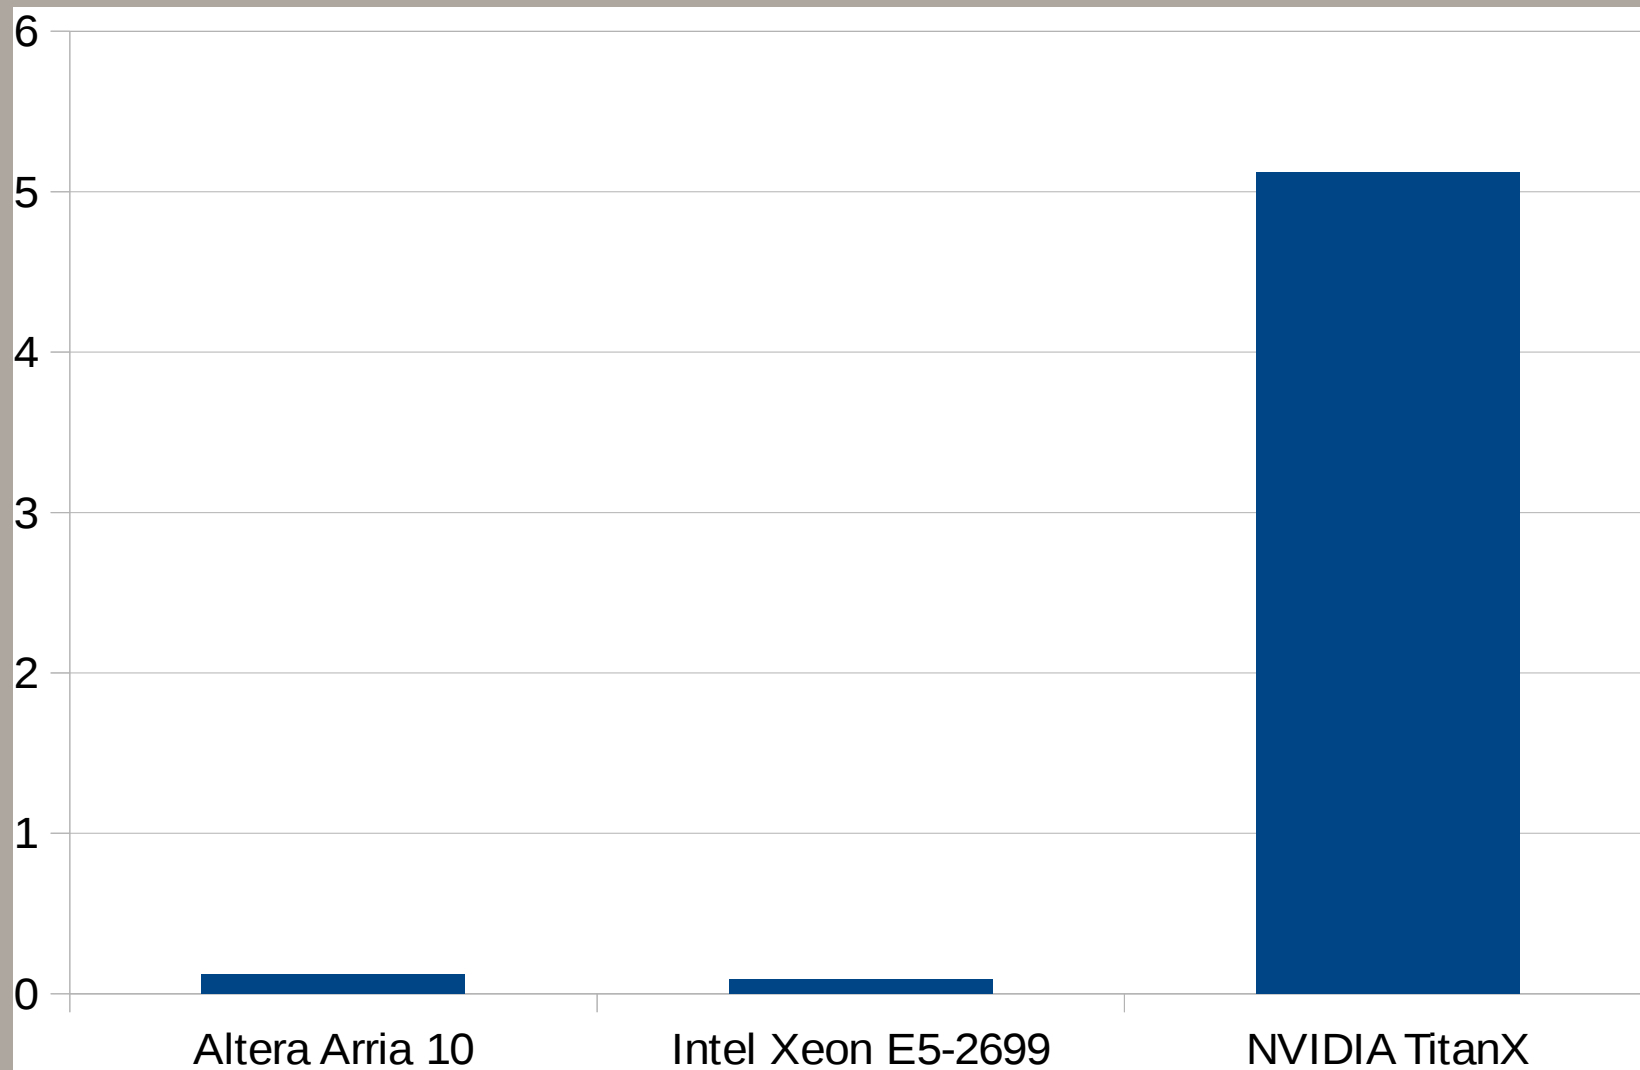


# AlexNet Images/Joule\*



\*Kudos to Ross Walker

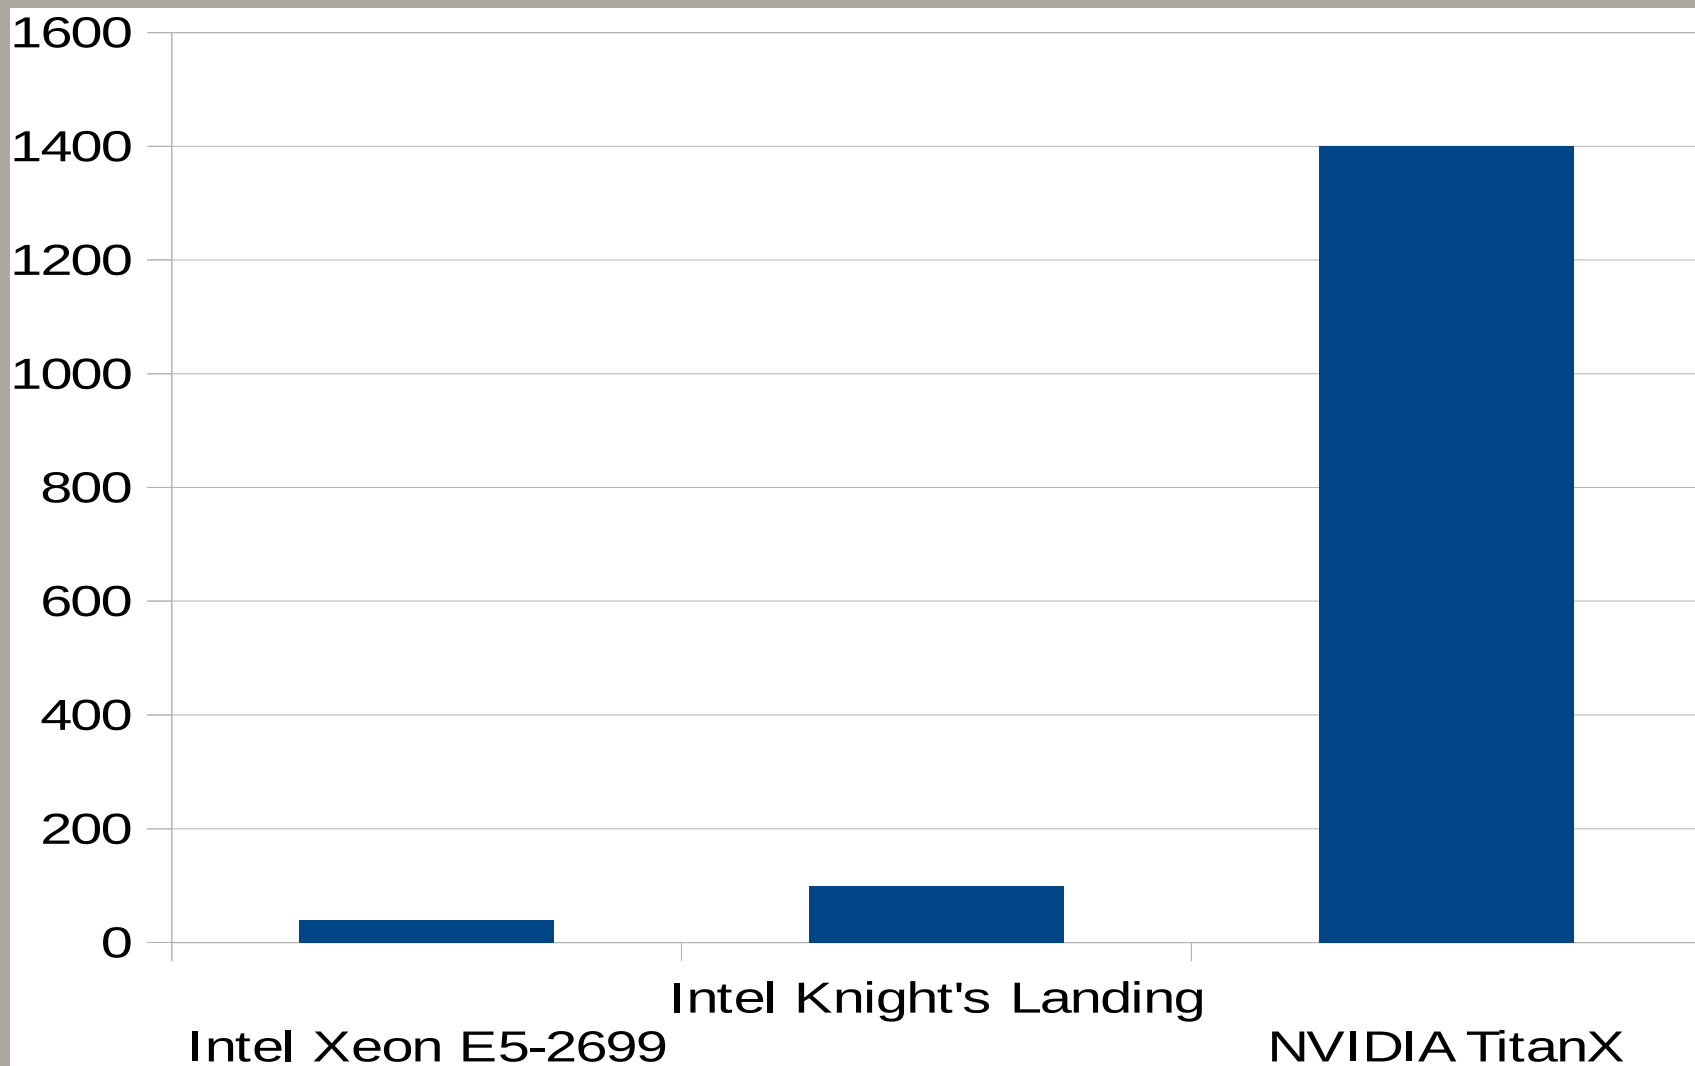
# AlexNet Images/s/\$



# What About Knight's Landing?

- Knight's Landing training performance projected from a HotChips talk (because Intel hates giving out real numbers unless they have to)...
- This is not good news for them, CPU training performance is awful...

# Projected KNL Training Performance

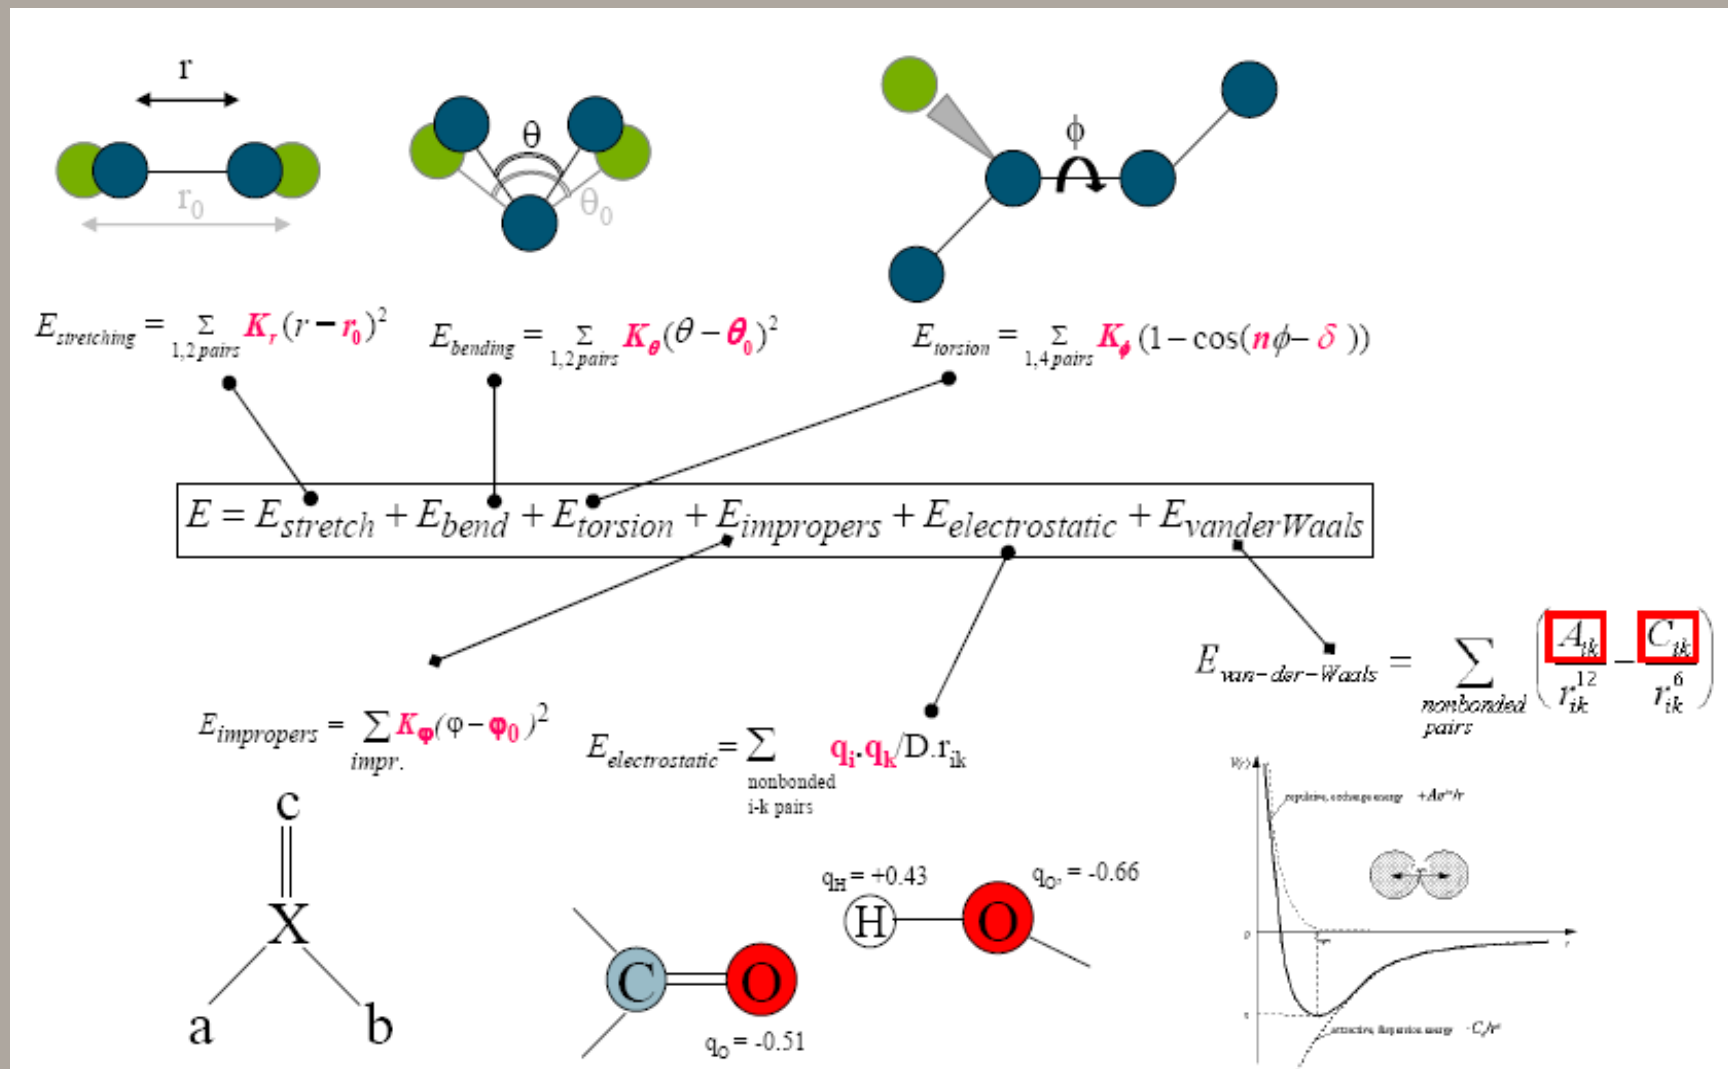




# Xeon Phi: A Trail of Tears

- KNL is ~6 TFLOPs, the HW can do a lot better
- But the engineers have been ordered to rely 100% on compiler improvements to implement “recompile and run”
- This is a fool's errand (IMO of course!)
- Nervana, NVIDIA and others have no such constraints
- Recompile and run is a no-win scenario
- Make OpenCL work across CPUs/Xeon Phi/FPGAs
- CUDA/OpenCL subsumes SIMD, multithreading, and multi-core

# AMBER Molecular Dynamics



# AMBER on GPUs

(or how to play a 30,720 string guitar)

On a CPU, the dominant performance spike is:

```
for (i = 0; i < N; i++)  
  for (j = i + 1; j < N; j++)  
    Calculate fij, fji;
```

$O(N^2)$  Calculation

If we naively ported this to a GPU, it would die the death of a thousand race conditions and memory overwrites


Solution: Reinvent mapreduce

# Mapreduced Molecular Dynamics


Force Matrix

j Atoms


i Atoms




Subdivide force matrix into 3 classes of independent tiles



Off-diagonal




On-diagonal




Redundant


# “Map” each nonredundant tile to a warp<sup>TM</sup>



Warp 0



Warp 1



Warp 2

⋮

⋮



Warp n

# Slow down, what's a warp?

The smallest unit of execution in a GPU  
similar to an AVX unit in a CPU

Up through GM2xx, it's groups of 32  
consecutive threads within the same core  
that execute in lockstep

GPU cores each run 8-64 warps at once  
on 4-6 vector units

May change in the future

Implements “lock-free computing”

# What's So Special About Warps?


`__shfl:` Exchanges data between warp threads

`__ballot:` Each bit gives state of a predicate for each warp thread

`__all:` True if predicate is true across all warp threads

`__any:` True if predicate is true on any warp thread

# What About The Reduce Part?



We've “mapped” the force matrix, now we have to “reduce” it to a force vector



# Two ways to Reduce

- Execute  $n$  separate  $n$ -way sums in parallel
  - Simple algorithm but it requires  $O(N^2)$  memory
- 
- Use Atomic Operations
  - No extra memory needed, but floating-point atomic operations are not deterministic

# Floating Point Math isn't Associative

$A + B == B + A$  (Commutative)

$A + B + C?$  (Associative)

$!= B + C + A$

$!= A + C + B$

$!= C + B + A$

So what? Big deal... Why should we care?

# Can you spot the broken GPU/Race Condition/Driver Bug/Thermal Issue/Software Bug?

GPU #1

ETot = -288,718.2326

ETot = -288,718,2325

GPU #2

ETot = -288,718.2326

Etot = -288,718,2326

# Let's make it easier...

GPU #1

GPU #2

$$ETot = -288,718.2326$$

$$ETot = -288,718.2326$$

$$ETot = -288,718,2325$$

$$Etot = -288,718,2326$$

# Non-Deterministic Accumulation

GPU #1

GPU #2

ETot = -288,456.6774

ETot = -288,458.5931

ETot = -288,453.8133

Etot = -288,454.1539

GeForce GPUs are not QAed for HPC, only gaming...

# Dynamic Range and Molecular Dynamics

32-bit floating point has approximately 7 significant figures

1.4567020  
+0.3046714

-----

1.7613730  
-1.4567020

-----

0.3046710  
Lost a sig fig

1456702.0000000  
+ 0.3046714

-----


1456702.0000000  
-1456702.0000000

-----

0.0000000  
Lost everything.

When it happens: PBC, SHAKE, and Force Accumulation in MD, backpropagation and recurrence in Neural Networks, esp. with FP16 gradients

# Dynamic Range Matters



# Deterministic Stable MD (using single-precision)

Acceptable force error is  $\sim 10^{-5}$  (as determined by D.E. Shaw)

Single-precision error is  $\sim 10^{-7}$

So calculate forces in single precision, but accumulate in extended precision

Before Kepler GPUs, we used double-precision and reduction buffers

GK104 (GTX 6xx) made it necessary to switch to 64-bit fixed point atomic Adds for accumulation because FP64 perf was reduced to 1/24 FP32



# **64-bit fixed point deterministic accumulation**

Each iteration of the main kernel in PMEMD uses 9 double-precision operations

Fermi double-precision was  $\frac{1}{4}$  to  $1/10^{\text{th}}$  of single-precision

GTX6xx double-precision is  $1/24^{\text{th}}$  single precision!

So accumulate forces in 64-bit fixed point

Fixed point forces are *\*perfectly\** conserved

3 double-precision operations per iteration

Integer extended math (add with carry) is 32-bit!

# Along Came GM2xx

- On GM2xx, double-precision (llrintf) was further reduced to 1/32 that of single-precision whilst nearly doubling attainable single-precision performance (GM200 versus GK110, GM204 versus GK104)
- Initially GM204 is slightly better than GTX 780, GM200 ~20% better than GK110
- Fortunately, we had a solution waiting in the wings that we developed for GK1xx

# Use 2 x FP32 (~48-bit FP)

Extended-Precision Floating-Point Numbers for GPU  
Computation - Andrew Thall, Alma College

[http://andrewthall.org/papers/df64\\_qf128.pdf](http://andrewthall.org/papers/df64_qf128.pdf)

High-Performance Quasi Double-Precision Method Using  
Single-Precision Hardware for Molecular Dynamics on GPUs –  
Tetsuo Narumi *et al.*

HPC Asia and APAN 2009

# Knuth & Dekker Summation

Represent ~FP48 as 2 floats

```
struct Accumulator {  
    float hs;  
    float ls;  
    Accumulator() : hs(0.0f), ls(0.0f) {}  
};
```

# Accumulation

```
void add_forces(Accumulator& a, float ys)
{
    float hs, ls, ws;

    // Knuth and Dekker addition
    hs      = a.hs + ys;
    ws      = hs - a.hs;
    a.hs    = ls;
    a.ls    = ys - ws;
}
```


# Conversion to 64-bit int

```
long long int upcast_forces(Accumulator& a)
{
    long long int l    = llrintf(a.hs * FORCESCALEF) +
                        llrintf(a.ls * FORCESCALEF);
    return l;
}
```

# NVIDIA fixes the problem

```
long long fast_llrintf(float x) {  
  
    float z = x * (float)0x1.00000p-32;  
    int hi = __float2int_rz( z );  
  
    float delta = x - ((float)0x1.00000p32*((float)hi));  
    int test = (__float_as_uint(delta) > 0xbf000000);  
    int lo = __float2uint_rn(fabsf(delta));  
    lo = (test) ? -lo: lo;  
    hi -= test;  
    long long res =  
        __double_as_longlong(__hilo_int2double(hi, lo));  
    return res;  
  
}
```

# AMBER Performance





# Summary

- Refactoring Molecular Dynamics into a mapreduce-like task decomposition has allowed performance to scale proportionally to GPU performance
- Refactoring for the next GPU generation is a 1-2 week task based on 7 years and 4 GPU generations
- Much less work than SSE/SSE2/SSE3/SSE4/AVX/AVX2/AVX512 hand-coded intrinsics (IMO of course)

# More AMBER?

Speed Without Compromise: Precision and  
Methodology/Innovation in the AMBER GPU  
MD Software

Ross Walker, April 7, 10:30 AM right here

# CPUs are looking more and more like GPUs

- CPU clocks haven't gone up in significantly in a decade
- Broadwell will have up to 22 physical cores and dual 8-way AVX2 units
- TitanX has 24 cores and 4 32-way vector units
- Later Skylake chips will have Dual AVX 512 units
- GPU-friendly algorithms are AVX-friendly algorithms

# Neural Networks\*

$$X_{L+1} = X_L * W_{L \rightarrow L+1}$$

$$\delta_L = \delta_{L+1} * W_{L \rightarrow L+1}$$


$$\Delta W = X_L^T * \delta_{L+1}$$

\*The definitive answer to whether you should take Calculus, Statistics and Linear Algebra in college

# Model Parallel Training

“My belief is that we’re not going to get human-level abilities until we have systems that have the same number of parameters in them as the brain.” - Geoffrey Hinton

# P2P Scatter/Gather Ops 2016\*




\*As seen (but implemented inefficiently) in the NVIDIA NCCL library

# P2P Ring Ops Performance\*

- AllReduce:  $2 * D * (N - 1) / N$
- Scatter/Gather/AllGather:  $D * (N - 1) / N$
- Reduce:  $D * (N - 1) / N$


\*NVLINK makes everything better, but we'll get to that...

# The AMBERnator (2013)



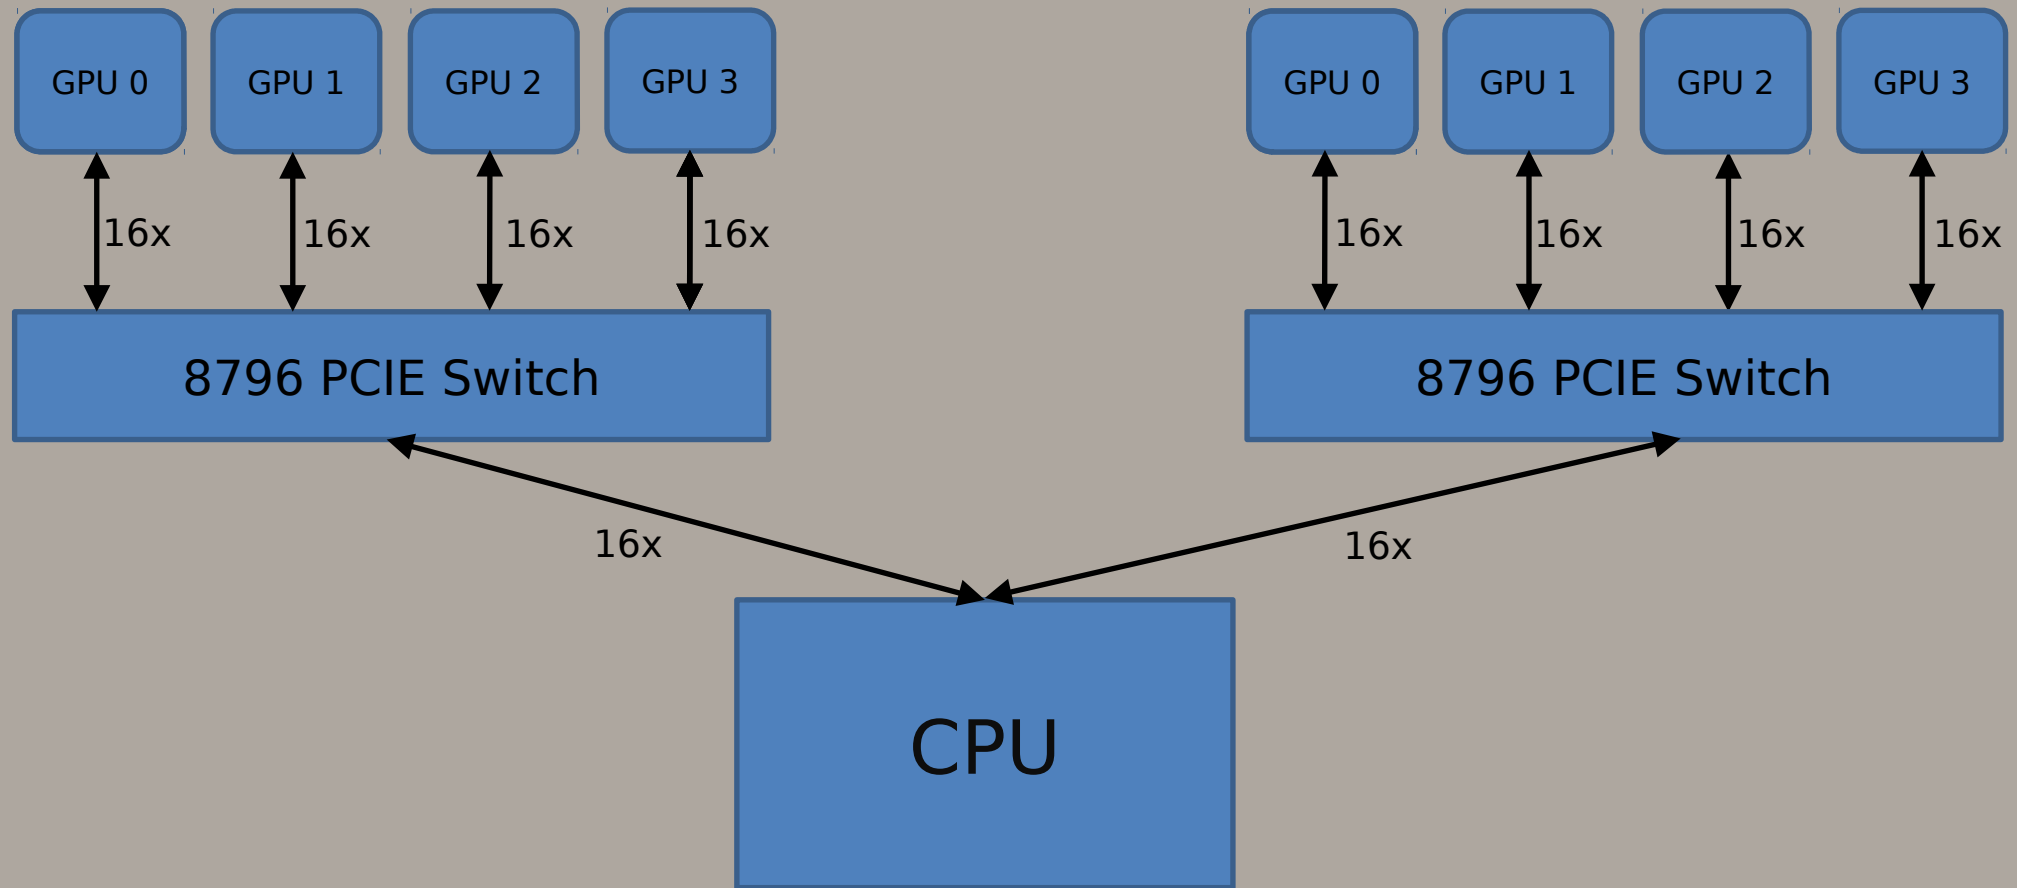


# Digits Dev Box (2015)\*



\*Maybe you can tell me the difference?


# Inefficient (2016)



# Intel hates P2P Bandwidth

	0	1	2	3	4	5	6	7
0	NA	25.03	25.02	25.01	15.97	15.97	14.73	15.97
1	25.03	NA	25.04	25.02	15.96	15.97	14.73	15.97
2	25.02	25.04	NA	25.02	15.97	15.96	14.73	15.96
3	25.02	25.03	25.02	NA	14.69	14.69	14.7	14.69
4	15.98	15.98	15.99	14.73	NA	25.02	25.04	25.03
5	15.98	15.98	15.98	14.73	25.03	NA	25.02	25.03
6	14.69	14.7	14.69	14.7	25.03	25.02	NA	25.03
7	15.98	15.97	15.98	14.73	25.04	25.04	25.03	NA


# Big Sur (Efficient, 2016)




# PLX loves P2P Bandwidth

	0	1	2	3	4	5	6	7
0	NA	24.97	24.96	24.95	24.95	24.95	24.96	24.95
1	24.97	NA	24.97	24.96	24.96	24.95	24.95	24.96
2	24.97	24.95	NA	24.95	24.96	24.96	24.95	24.95
3	24.95	24.95	24.95	NA	24.94	24.96	24.96	24.96
4	24.95	24.95	24.95	24.95	NA	24.94	24.95	24.94
5	24.95	24.95	24.94	24.94	24.95	NA	24.94	24.95
6	24.95	24.95	24.95	24.94	24.94	24.94	NA	24.95
7	24.94	24.94	24.95	24.94	24.95	24.95	24.96	NA


# P2P Ring Implementation



# P2P Ring Simplified




# Model Parallel Data






# Model Parallel Weights



\*Have you spotted the kryptonite yet?

# Other Model Parallel Weights\*





\*Weight Subdivision Style Matters!

# Other Other Model Parallel


- Layer by layer subdivision of network
- Sir not appearing in this talk
- Supported by Tensorflow

# One Weird(er) Trick\*




\* Perform N SGEMM operations and reduce the outputs over N-1 communication steps if the model outputs are smaller than the model inputs

# Reduction Flowchart



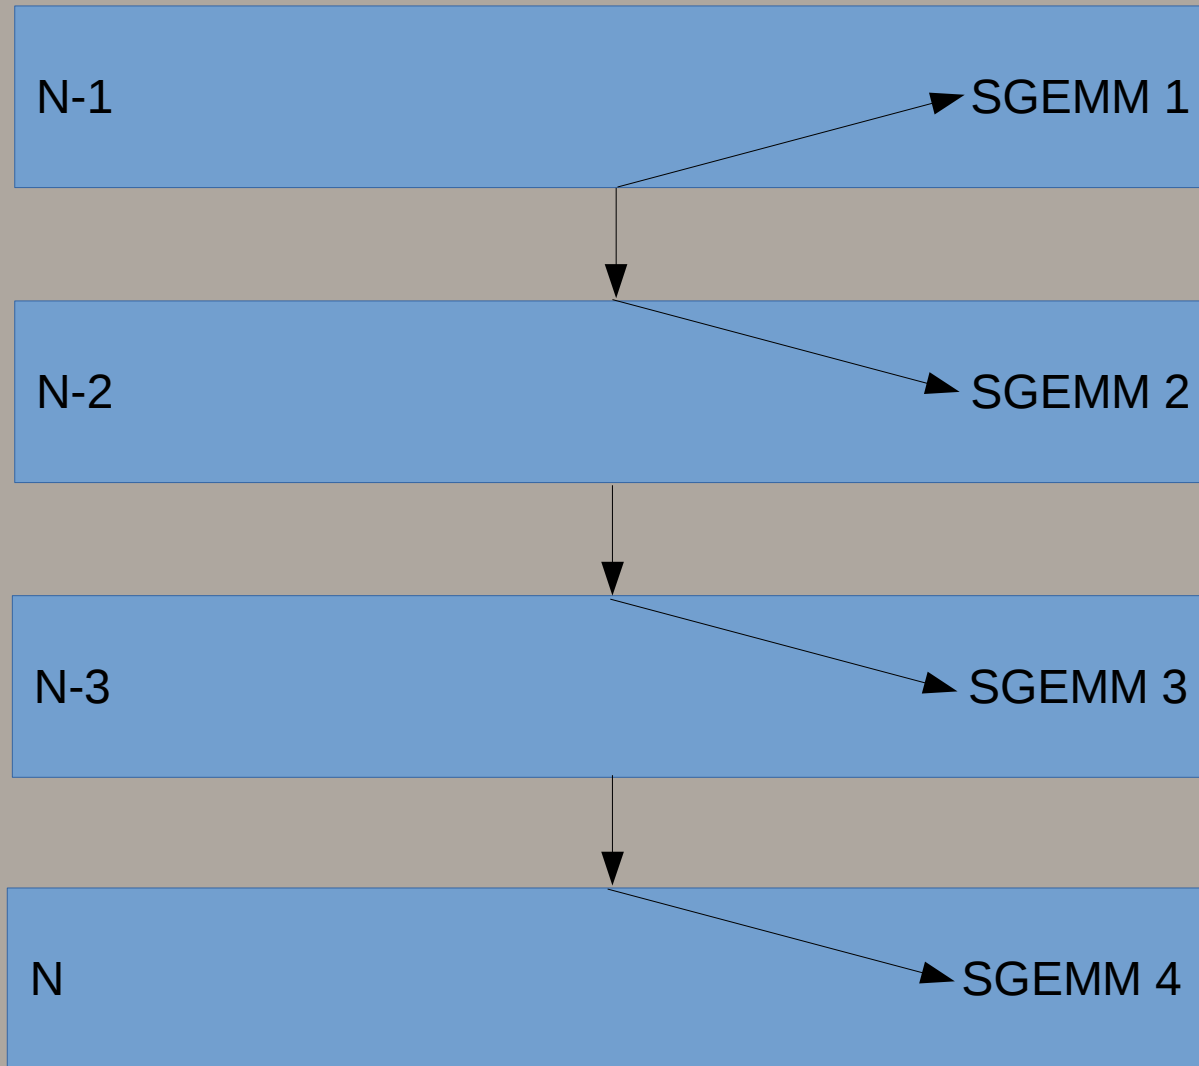
# And the other Weirder Trick\*

$$\begin{matrix} X_{L:1} \\ X_{L:2} \\ X_{L:3} \\ X_{L:4} \end{matrix} \times \begin{matrix} W_{L \rightarrow L+1:1} \end{matrix} = \begin{matrix} X_{L+1:1} \end{matrix}$$



\*Scatter the inputs over  $N-1$  communication steps and SGEMMs if the model inputs are smaller than the model outputs

# Gather Flowchart



# Overlappppping Computation/Communication

- TitanX is ~6.6 TFLOPS
- PCIE BW is ~12.5 GB/s if you don't buy crap HW
- $6.6 \text{ TFLOPS} / 3.125 \text{ Gfloats} \approx 2000 \text{ FLOPS}$
- So if you have ~1000 FMADs per output per GPU, you can run at SOL\*

\*This requires efficient SGEMM, cuBLAS is fraught...




# TitanX is ~6.6 TFLOPS, cuBLAS is whatever it feels like...


- For small batch sizes, cuBlas is anywhere from 1/10th to 1/6th of SOL
- cuBLAS **hates hates hates** narrow matrices
- cuBLAS is obsessed with multiples of 128
- Scott Gray's BLAS kernels have no such psychological issues

<https://github.com/NervanaSystems/neon>


# Fast Case (many outputs)



# Slow Case (few outputs)



# Solution: Subdivide each SGEMM within each GPU\*



\*Independently discovered by Scott Gray

# Do you need a DGX1?

- 85 TFLOPS FP32 (10.6 TFLOPS per GPU) no FP16 for now
- 20 GB/s channels connected in a cube ( $N == 8$ )

Reduction:  $2 * D / N$  vs  $\sim 1.6 * D * (N - 1) / N$

Gather:  $2 * D / N$  vs  $\sim 1.6 D * (N - 1) / N$

AllReduce:  $0.5 * D$  vs  $\sim 3.2 * D * (N - 1) / N$

Significant reduction in communication costs, but is AlexNet communication-limited?

# Are you data-parallel?

- AlexNet has ~61M parameters
- We'll assume a batch size of 128 and Soumith Chintala's training perf numbers for TitanX scaled up by ~1.6 to arrive at 2,884 images/s FP32
- 16 images at 2,884 images/s is ~5.5 ms
- AllReducing 61M (244 MB) parameters at 20 GB/s is ~6 ms (buried 5.5 ms of backprop for overlapping copy and compute) for a final result of 0.4 ms or nearly free.
- Using P2P, this would take ~34 ms, \$129K is a bargain!

# Alex Krizhevsky to the Rescue! (or are you model-parallel?)

- AlexNet has ~61M parameters. ~4.3M of which are convolutional (data-parallel) and ~56.7M of which are fully-connected (model-parallel)
- Fully connected layers at a batch size of 128 is ~1.7M neurons
- P2P allReduce of 4.3M parameters takes ~2.4 ms
- P2P gather/reduction of 1.7M neurons is ~0.5 ms
- 2.9 ms is  $\ll$  5.5 ms so once again it's free(tm)
- It's also faster than NVLINK data-parallel...
- NVLINK model-parallel would of course win...

TLDR: Go Model  
Parallel or Go  
Home...



# Summary

- GPUs still rule HPC/Machine Learning
- Rethink algorithms into parallel-friendly implementations instead of waiting for compilers to do this for you (because they won't)
- Who needs DGX1 if we utilize model parallelism?

# Acknowledgments (Amazon)

Matias Benitez

Kiuk Chung

Leo Dirac

Rejith Joseph George

Mitchell Goodman

Sebastian Gunningham

Shruti Kamath

Oleg Rybakov

Srikanth Thirumulai

Jane You

# Acknowledgments (AMBER)

David Case

Romelia Salomon-Ferrer

Ben Madej

Perri Needham

Levi Pierce

Adrian Roitberg

Jason Swails

Ross Walker

# Acknowledgments (NVIDIA)

Jonathan Bentz

Mark Berger

Jerry Chen

Kate Clark

Simon Layton

Duncan Poole

Sarah Tariq